

Course Logistics

- CS533: Intelligent Agents and Decision Making
 - ▲ M, W, F: 1:00—1:50
 - ▲ Instructor: Alan Fern (KEC2071)
 - ▲ Office hours: Th, F 10-11
 - ▲ Course website (link on instructor's home page) has
 - Lecture notes, Assignments, and HW Solutions
- The course will cover basics of AI planning with some advanced topics
 - ▲ Prepare you to formulate planning problems, apply algorithms, and read research papers on planning
- Homework:
 - ▲ Assigned and collected regularly
 - ▲ Not graded for correctness, but rather for “completion with good effort”
 - ▲ Must “complete with good effort” 90% of homework or a letter grade will be deducted from final grad
- Grade based on:
 - ▲ 20% Midterm exam (in class)
 - ▲ 30% Final exam (take home)
 - ▲ 30% Final Project (during last month)
 - ▲ 20% 2 mini-projects

Course Outline

- Classical planning
 - ▲ Basic algorithmic paradigms
- Markov Decision Processes (MDPs)
 - ▲ Basic definitions and solution techniques
- Reinforcement learning
 - ▲ Basic algorithms
- Monte-Carlo Planning
 - ▲ Basic algorithms

Homework 1

- Due January 12
- Download from course web-page

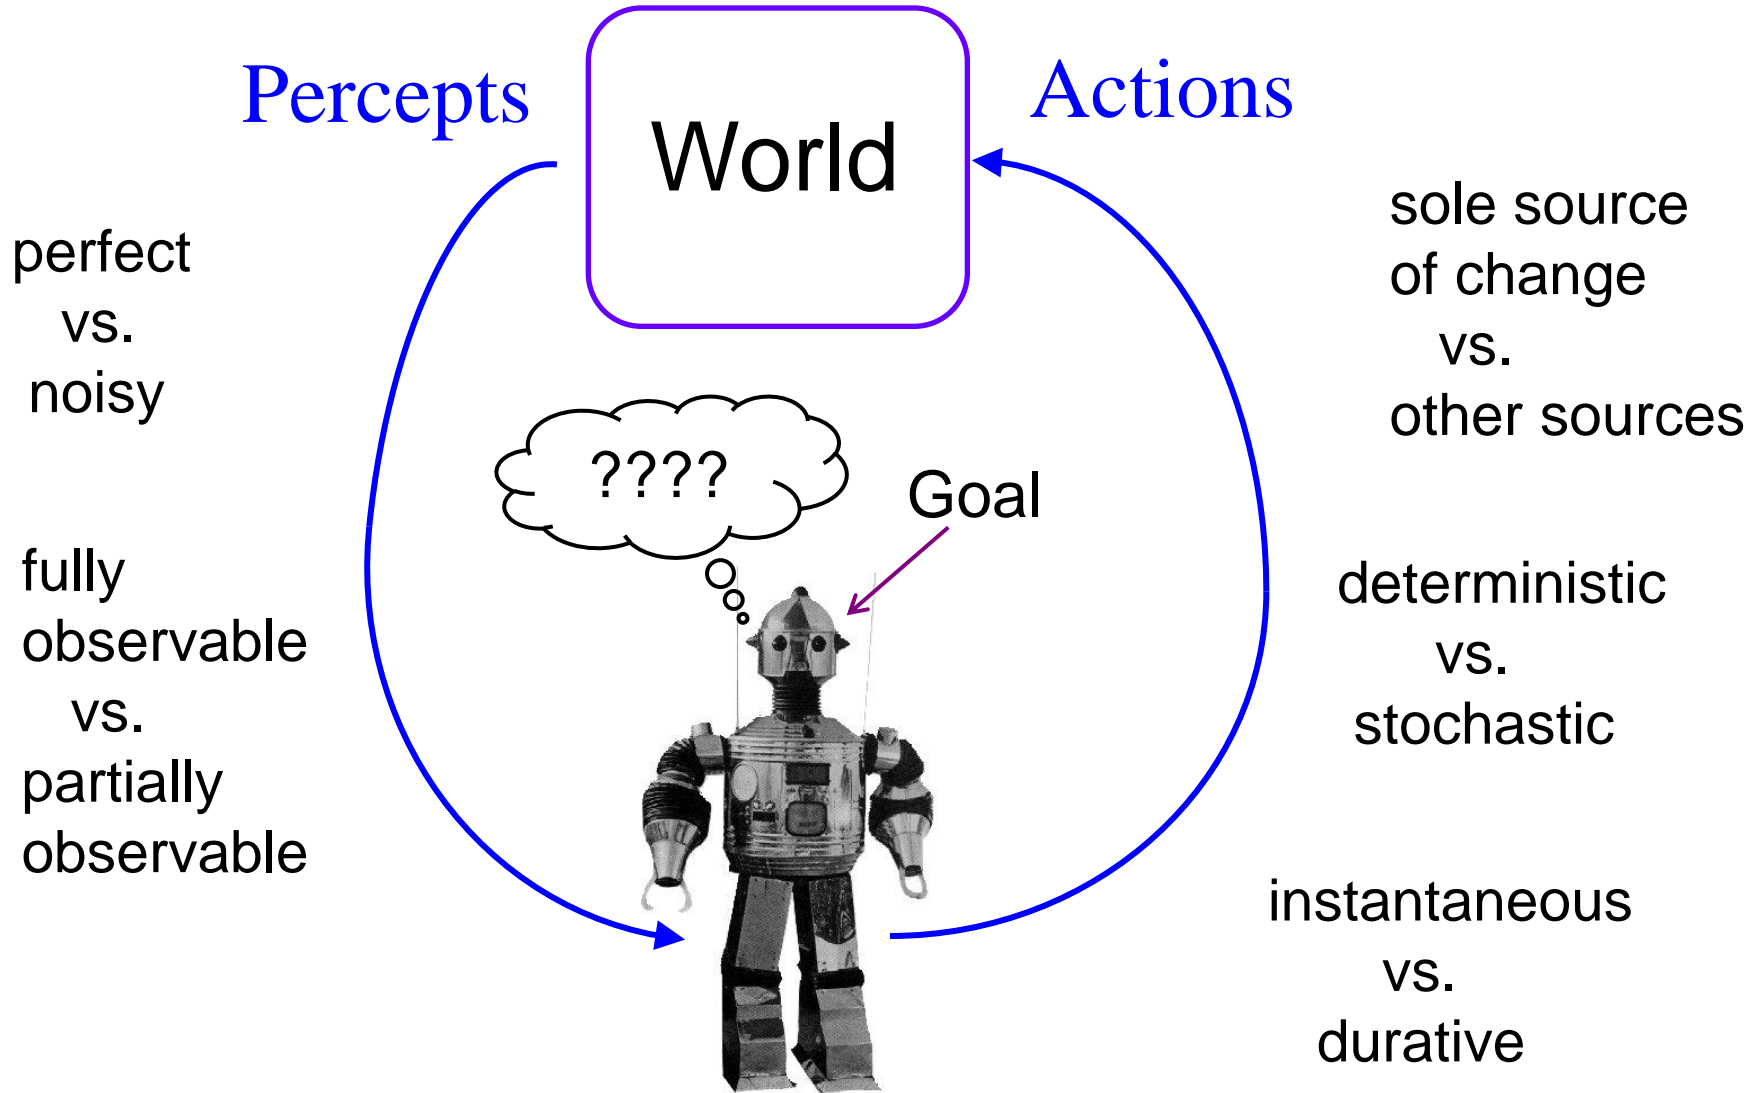
Classical STRIPS Planning

Alan Fern *

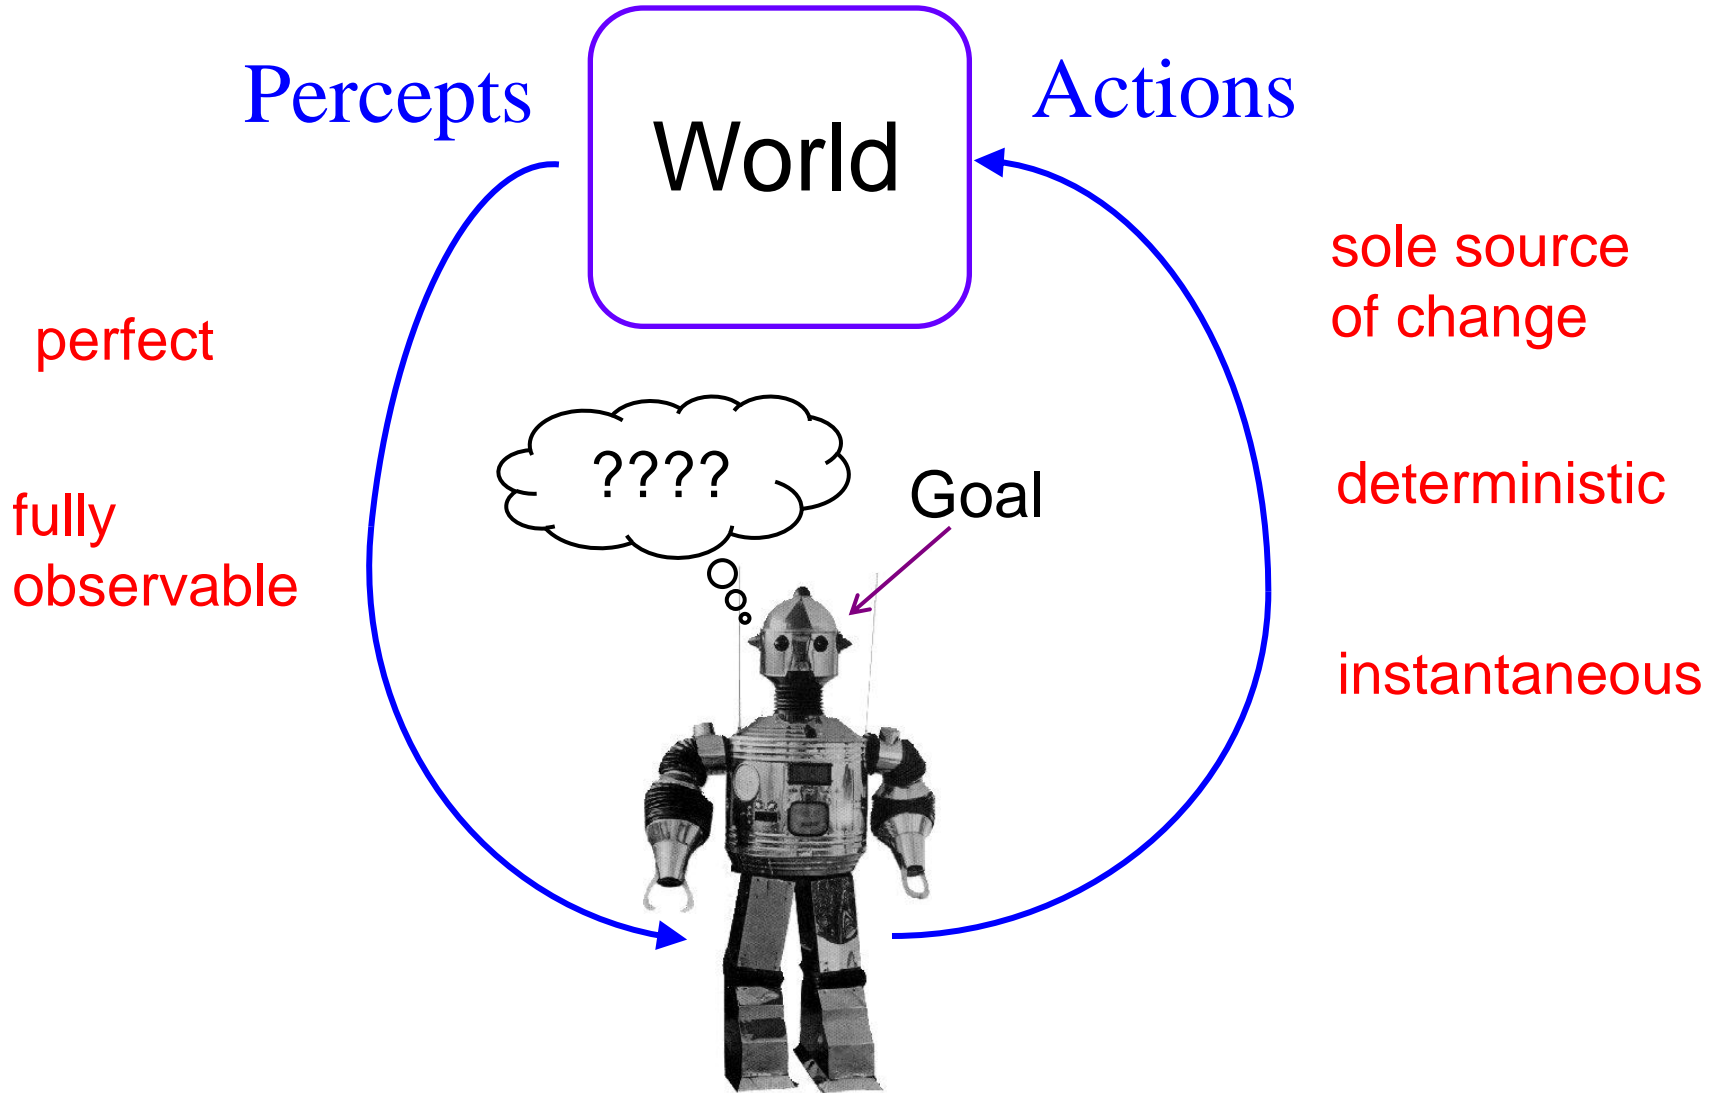
- Classical planning assumptions
- STRIPS action language
- What is a planning problem?
 - ▲ complexity of planning

* Based in part on slides by Daniel Weld.

AI Planning



Classical Planning Assumptions



Why care about classical planning?

- Most advances seen first in classical planning
- Many stabilized environments ~satisfy classical assumptions (e.g. robotic crate mover)
 - ▲ It is possible to handle minor assumption violations through replanning and execution monitoring


“ This form of solution has the advantage of relying on widely-used (and often very efficient) classical planning technology” Boutilier, 2000

- Ideas from classical planning techniques often form the basis for developing non-classical planning techniques
 - ▲ E.g. recent work uses classical planners for probabilistic planning [Yoon et. al. 2008]

Representing States

World states are represented as sets of facts.

We will also refer to facts as propositions.




holding(A)
clear(B)
on(B,C)
onTable(C)

State 1

handEmpty
clear(A)
on(A,B)
on(B,C)
onTable(C)

State 2



Closed World Assumption (CWA):


Fact not listed in a state are assumed to be false. Under CWA we are assuming the agent has full observability.

Representing Goals


Goals are also represented as sets of facts.

For example $\{ \text{on(A,B)} \}$ is a goal in the blocks world.

A **goal state** is any state that contains all the goal facts.



State 1

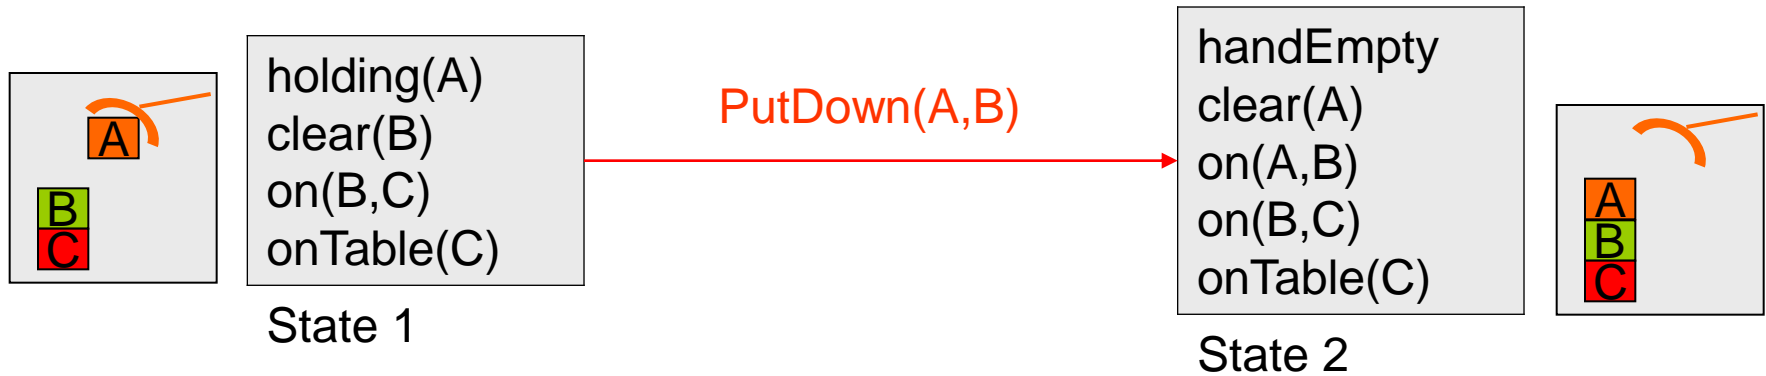


State 2

State 1 is a goal state for the goal $\{ \text{on(A,B)} \}$.

State 2 is not a goal state.

Representing Action in STRIPS



A STRIPS action definition specifies:

- 1) a set PRE of preconditions facts
- 2) a set ADD of add effect facts
- 3) a set DEL of delete effect facts

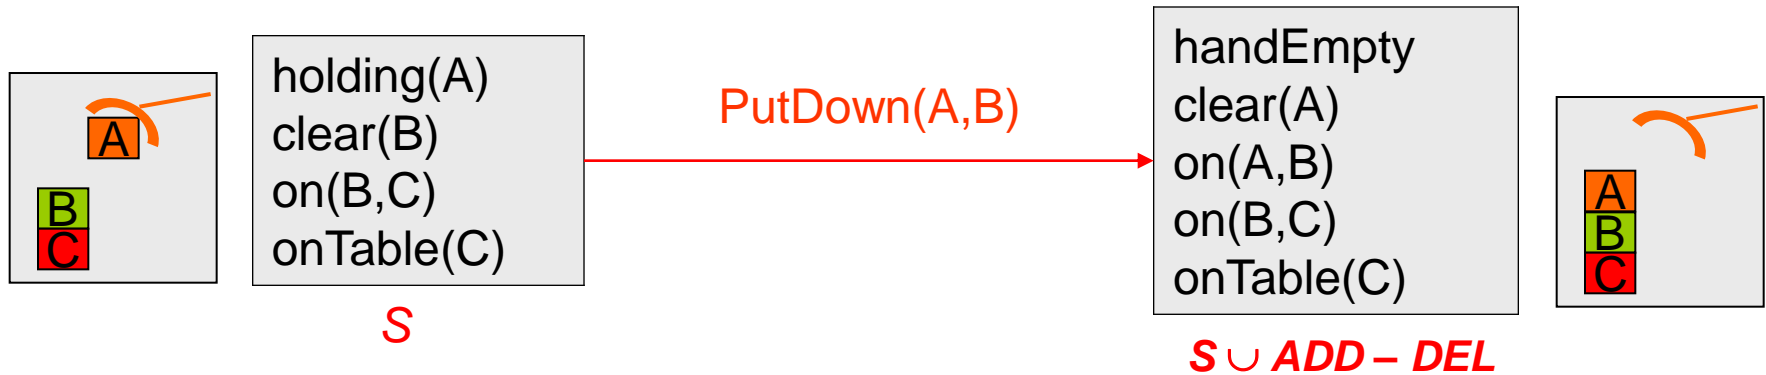
PutDown(A,B):

PRE: { holding(A), clear(B) }

ADD: { on(A,B), handEmpty, clear(A) }

DEL: { holding(A), clear(B) }

Semantics of STRIPS Actions



- A STRIPS action is **applicable** (or allowed) in a state when its preconditions are contained in the state.
- Taking an action in a state **S** results in a new state **$S \cup \text{ADD} - \text{DEL}$** (i.e. add the add effects and remove the delete effects)

PutDown(A,B):

PRE: { `holding(A)`, `clear(B)` }

ADD: { `on(A,B)`, `handEmpty`, `clear(A)` }

DEL: { `holding(A)`, `clear(B)` }

STRIPS Planning Problems


A **STRIPS planning problem** specifies:

- 1) an initial state S
- 2) a goal G
- 3) a set of STRIPS actions

Objective: find a “short” action sequence reaching a goal state, or report that the goal is unachievable

Example Problem:

Solution: (**PutDown(A,B)**)



holding(A)
clear(B)
onTable(B)

Initial State

on(A,B)

Goal

PutDown(A,B):

PRE: { holding(A), clear(B) }

ADD: { on(A,B), handEmpty, clear(A) }

DEL: { holding(A), clear(B) }

PutDown(B,A):

PRE: { holding(B), clear(A) }


ADD: { on(B,A), handEmpty, clear(B) }

DEL: { holding(B), clear(A) }

STRIPS Actions

Propositional Planners

- For clarity we have written propositions such as `on(A,B)` in terms of objects (e.g. A and B) and predicates (e.g. `on`).
- However, the planners we will consider ignore the internal structure of propositions such as `on(A,B)`.
- Such planners are called **propositional planners** as opposed to first-order or relational planners
- Thus it will make no difference to the planner if we replace every occurrence of “`on(A,B)`” in a problem with “`prop1`” (and so on for other propositions)
- It feels wrong to ignore the existence of objects. But currently propositional planners are the state-of-the-art.



STRIPS Action Schemas

For convenience we typically specify problems via action schemas rather than writing out individual STRIPS action.

Action Schema: (x and y are variables)

PutDown(x,y):

PRE: { holding(x), clear(y) }
ADD: { on(x,y), handEmpty, clear(x) }
DEL: { holding(x), clear(y) }

PutDown(B,A):

PRE: { holding(B), clear(A) }
ADD: { on(B,A), handEmpty, clear(B) }
DEL: { holding(B), clear(A) }

■ ■ ■ ■

PutDown(A,B):

PRE: { holding(A), clear(B) }
ADD: { on(A,B), handEmpty, clear(A) }
DEL: { holding(A), clear(B) }

- Each way of replacing variables with objects from the initial state and goal yields a “ground” STRIPS action.
- Given a set of schemas, an initial state, and a goal, propositional planners compile schemas into actions and then ignore the existence of objects thereafter.

STRIPS Versus PDDL

- Your book refers to the PDDL language for defining planning problems rather than STRIPS
- The **Planning Domain Description Language (PDDL)** was defined by planning researchers as a standard language for defining planning problems
 - ▲ Includes STRIPS as special case along with more advanced features
 - ▲ Some simple additional features include: type specification for objects, negated preconditions, conditional add/del effects
 - ▲ Some more advanced features include allowing numeric variables and durative actions
- Most planners you can download take PDDL as input
 - ▲ Majority only support the simple PDDL features (essentially STRIPS)
 - ▲ PDDL syntax is easy to learn by look at examples packaged with planners, but a definition of the STRIPS fragment can be found at:
<http://eecs.oregonstate.edu/ipc-learn/documents/strips-pddl-subset.pdf>

Properties of Planners

- A planner is **sound** if any action sequence it returns is a true solution
- A planner is **complete** if it outputs an action sequence or “no solution” for any input problem
- A planner is **optimal** if it always returns the shortest possible solution

Is optimality an important requirement?


Is it a reasonable requirement?

Planning as Graph Search

- It is easy to view planning as a graph search problem
- Nodes/vertices = possible states
- Directed Arcs = STRIPS actions
- Solution: path from the initial state (i.e. vertex) to one state/vertices that satisfies the goal

Search Space: Blocks World

Graph is finite



Planning as Graph Search

- Planning is just finding a path in a graph
 - ▲ Why not just use standard graph algorithms for finding paths?
- **Answer:** graphs are exponentially large in the problem encoding size (i.e. size of STRIPS problems).
 - ▲ But, standard algorithms are poly-time in graph size
 - ▲ So standard algorithms would require exponential time
- Can we do better than this?

Complexity of STRIPS Planning

PlanSAT

Given: a STRIPS planning problem

Output: “yes” if problem is solvable, otherwise “no”


- PlanSAT is decidable.
 - ▲ Why?
- In general PlanSAT is PSPACE-complete!
Just finding a plan is hard in the worst case.
 - ▲ even when actions limited to just 2 preconditions and 2 effects

Does this mean that we should give up on AI planning?

NOTE: PSPACE is set of all problems that are decidable in polynomial space.
PSPACE-complete is widely believed to strictly contain NP.

Satisficing vs. Optimality

- While just finding a plan is hard in the worst case, for many planning domains, finding a plan is easy.
- However finding optimal solutions can still be hard in those domains.
 - ▲ For example, optimal planning in the blocks world is NP-complete.
- In practice it is often sufficient to find “good” solutions “quickly” although they may not be optimal.
 - ▲ This is often referred to as the “satisficing” objective.
 - ▲ For example, producing approx. optimal blocks world solutions can be done in linear time. How?



Satisficing

- Still finding satisficing plans for arbitrary STRIPS problems is not easy.
 - ▲ Must still deal with the exponential size of the underlying state spaces
- Why might we be able to do better than generic graph algorithms?
- **Answer:** we have the compact and structured STRIPS description of problems
 - ▲ Try to leverage structure in these descriptions to intelligently search for solutions
- We will now consider several frameworks for doing this, in historical order.